

Milan Mossé
milan_mosse@stanford.edu
mmosse19.github.io
+1 (626) 710-2177

Employment

Embedded Ethics Fellow, Stanford University 2026-current
The McCoy Family Center for Ethics in Society &
The Institute for Human-Centered Artificial Intelligence

Areas

Specialization: Ethics and Political Philosophy (including AI)
Competence: Philosophy of Action, Language, Race, and Science; Logic

Education

Ph.D., Philosophy, University of California, Berkeley 2026
Committee: Johann Frick (co-chair), R. Jay Wallace (co-chair), Anna Stilz
M.S., Computer Science, Stanford University 2024
Research supervisors: Thomas Icard, Omer Reingold, Li-Yang Tan
B.A., Philosophy (Honors) & B.S., Mathematics, Stanford University 2020
Thesis advisors: Michael Bratman, Barry Maguire

Peer-reviewed publications

 [Google Scholar](#)  [PhilPapers](#)

How Requests Create Reasons 2026
Philosophy & Phenomenological Research
Stable Voting and the Splitting of Cycles 2026
Association for the Advancement of Artificial Intelligence
Wesley H. Holliday, Milan Mossé, Chase Norman, Eric Pacuit, Cynthia Wang
How to Count Sore Throats 2025
Analysis
Léa Bourguignon, Milan Mossé
On Probabilistic and Causal Reasoning with Summation Operators 2025
Journal of Logic and Computation
Duligur Ibeling, Thomas Icard, Milan Mossé

Social Choice for AI Alignment: Dealing with Diverse Human Feedback	2024
<i>International Conference on Machine Learning</i>	
Vincent Conitzer, Rachel Freedman, Jobst Heitzig, Wesley Holliday, Bob Jacobs, Nathan Lambert, Milan Mossé, Eric Pacuit, Stuart Russell, Hailey Schoelkopf, Emanuel Tewolde, William Zwicker	
Probing the Quantitative-Qualitative Divide in Probabilistic Reasoning	2024
<i>Annals of Pure and Applied Logic</i>	
Duligur Ibeling, Thomas Icard, Krzysztof Mierzewski, Milan Mossé	
Is Causal Reasoning Harder than Probabilistic Reasoning?	2024
<i>Review of Symbolic Logic</i>	
Milan Mossé, Duligur Ibeling, Thomas Icard	
Multiplicative Metric Fairness Under Composition	2023
<i>Foundations of Responsible Computing</i>	
A Generalization of the Satisfiability Coding Lemma and Its Applications	2022
<i>International Conference on Theory and Applications of Satisfiability Testing</i>	
Milan Mossé, Harry Sha, Li-Yang Tan (Best Theory Paper Award)	
Conditional Negative Sampling for Contrastive Learning of Visual Representations	2021
<i>International Conference on Machine Learning</i>	
Mike Wu, Milan Mossé, Chengxu Zhuang, Daniel Yamins, Noah Goodman	
Zero Shot Learning for Code Education: Rubric Sampling with Deep Learning Inference	2019
<i>Association for the Advancement of Artificial Intelligence</i>	
Mike Wu, Milan Mossé, Noah Goodman, Chris Piech (Outstanding Student Paper Award)	

In the pipeline

Why Not Ask?

R&R, *Ethics*

Modeling Discrimination with Causal Abstraction

Milan Mossé, Kara Schechtman, Frederick Eberhardt, Thomas Icard

R&R, *Philosophers' Imprint*

[Nominated for Best Paper, FAccT 2026](#)

Individualism and Its Limits

Léa Bourguignon, Jonas Harney, Milan Mossé

R&R, *Noûs*

Giving Grace

Invited to submit to *Oxford Studies in Normative Ethics*

Why AI Alignment Needs Moral Foundations

Invited to submit to *Philosophy & Technology*

Edited volume

Limited Aggregation, *Philosophical Studies* 2028 (expected)
Guest editor, with Léa Bourguignon & Jonas Harney

Awards & fellowships

Graduate Fellow, Kavli Center for Ethics, Science, and the Public, UC Berkeley	2024-25
Dissertation Completion Fellowship, UC Berkeley	2025
Outstanding GSI Award (awarded to fewer than 10% of TAs), UC Berkeley	2023
Cota Robles Fellowship, UC Berkeley	2020
Major Grant, Stanford	2019
Mellon Mays Undergraduate Fellow, Stanford	2017-29
Chappell-Lougee Scholarship, Stanford	2017

Selected presentations

Betting on Misfortune (with Johann Frick) The Normative Foundations of the Market, University of Pennsylvania*	2027
Corruption as the Devaluation of Membership (with Eliza Wells) Central APA Colloquium	2027
Giving Grace Arizona Workshop in Normative Ethics	2027
Philosophy Forum, UC Berkeley*	2026
How Questions Structure Plans (with Seth Yalcin; various titles) 5th Annual Conference of the Philosophy of Language Association, UCLA	2026
Planning Workshop, Stanford University*	2026
Eastern APA Symposium*	2026
Agency and Intentions in Language, Harvey Mudd College	2024
Individualism and Its Limits Rocky Mountain Ethics Congress, University of Colorado Boulder	2026
Choice Group, London School of Economics*	2026
Workshop on Moral Aggregation, University of Manchester	2025
Modeling Discrimination with Causal Abstraction (with Kara Schechtman) To Be or Not to Be Included in a Causal Model, University of Pittsburgh	2026
Departmental Talk, Emory University*	2025
CEPDISC, Aarhus University	2024
Stanford Logic Seminar*	2024
How Requests Create Reasons	

Berkeley-London Conference	2026
Eastern APA Symposium	2026
Reasoning about Confirmation	
Probability Logic Workshop, Carnegie Mellon University	2025
Why Not Ask? (various titles)	
Central APA Symposium	2025
Rocky Mountain Ethics Congress, University of Colorado Boulder	2024
*invited; all others via application	

Teaching

University of California, Berkeley

TA, Philosophy of Science, with Ezra Rubenstein	Fall 2025
TA, Political Philosophy: Equality and Identity, with Véronique Munoz-Dardé	Fall 2024
Co-Instructor, Individual Morality and Social Justice, with Russell McIntosh	Summer 2024
TA, Political Philosophy, with Véronique Munoz-Dardé	Fall 2023
TA, Contemporary Ethical Issues, with Tim Crockett	Summer 2023
TA, Bioethics, with Johann Frick	Fall 2022
TA, Philosophy of Race, Ethnicity, and Citizenship, with Tim Crockett	Summer 2022
TA, Individual Morality and Social Justice, with Johann Frick	Spring 2022
TA, Nature of Mind, with John Campbell	Fall 2021

Stanford University

TA, Incentives in Computer Science, with Aviad Rubenstein	Spring 2024
TA, Design and Analysis of Algorithms, with Karey Shi	Summer 2020
TA, Design and Analysis of Algorithms, with Mary Wootters	Winter 2020
TA, Programming Abstractions, with Marty Stepp	Winter 2018
TA, Programming Methodology, with Mehran Sahami	Fall 2018

Service

Diversity, equity, and inclusion

Graduate Representative, Equity Task Force, Focus on Trans Inclusion	2021-22, 2024-25
Co-Organizer, Minorities and Philosophy	2021-22, 2023-24
Co-Organizer, COMPASS Workshop	Fall 2022
Un-Naming Proposal Co-Author, Moses Hall Name Review Committee	Spring 2021

Other organizing, chairing, and refereeing

Organizer, Job Market Materials Workshop	Summer 2025, 2026
Co-Organizer, MANCEPT Limited Aggregation Workshop	Fall 2025
Founder, Co-Organizer, & Mentor, Directed Reading Program	Spring 2022, 2023, 2025
Organizer, Graduate Work-in-Progress Talks	2020-21, 2023-2024
Grad Organizer, Social Choice for AI Ethics & Safety Workshop	Spring 2023
Grad Organizer, Department Colloquia	2022-23
Organizer, Faculty Work-in-Progress Talks	Spring 2021/2023, 2023-24
Organizer, Latin American Philosophy Reading Group	Spring 2022
Organizer, Feminism and Masculinity Reading Group	2021-22

Referee: Computational Complexity Conference; Conference on Neural Information Processing Systems (×9); International Conference on Machine Learning (×5); International Colloquium on Automata, Languages and Programming; International Symposium on Algorithms and Computation; Journal of Logic and Computation; Journal of Moral Philosophy (×5); Logic Journal of the IGPL (×3); Noûs; Philosophical Studies (×2); Philosophy; Philosophy & Technology; Review of Symbolic Logic; Symposium on Discrete Algorithms; Thought; Utilitas

Skills

Languages: Spanish (professional proficiency) and German (reading) *Coding:* C, C++, Python

References

Johann Frick

Associate Professor of Philosophy
University of California, Berkeley
jfrick@berkeley.edu

Thomas F. Icard

C. I. Lewis Professor of Philosophy and (by courtesy) Professor of Computer Science
Stanford University
icard@stanford.edu

Véronique Munoz-Dardé (teaching reference)

Professor of Philosophy
University College London
Mills Adjunct Professor of Philosophy
University of California, Berkeley
munoz_darde@berkeley.edu

R. Jay Wallace

William & Trudy Ausfahl Chair in Philosophy
University of California, Berkeley
rjw@berkeley.edu

Seth Yalcin

Professor of Philosophy
University of California, Berkeley
yalcin@berkeley.edu

Reasons by Request

Dissertation Summary

You are not required to stop and aid every stranger you encounter. Seeing a lost tourist, you can permissibly decide not to help, thinking to yourself, “I’ll help the next time I see someone in their shoes.” But this thought rings hollow when the lost stranger *requests* your help. Your normative situation has changed. You must now cite a reason for not helping or else provide assistance. Requesters wield what is, on reflection, a remarkable power: to create reasons at will. The central puzzle of requests is to explain how this is possible. How can requesters give us reasons to do what they want, even when they stand in no relation of authority over us? And what would be lost if we saw requests as merely reminding us of reasons that exist already, rather than creating new ones?

My dissertation begins by explaining how requests create reasons. At the center of my account is the idea that we have reasons to grant requests when and because doing so improves our compliance with our duties. For example, we have a duty of mutual aid, but because this duty gives us the discretion to decide exactly when and how we will help others, we are liable to fall short of it, by repeatedly postponing compliance. When a lost stranger asks you for help, their request gives you a reason *here and now* to help *this person*. The concrete nature of the request makes it harder to keep postponing your duty of mutual aid. In general, requests create reasons because they thereby help us to comply better with our flexible moral duties.

I next consider how the values served by the power to request constrain its permissible use: What do we owe to the recipients of our requests? I approach this question by way of an asymmetry between requests and offers. It is often not okay to request an act, but perfectly fine to accept an offer for it. What might explain this? After all, the request and offer result in the same duty upon their acceptance, and they concern the same act. I argue that some requests are wrongful *precisely because they create reasons*. Creating a reason often generates a double-bind for the requestee: if they act as requested, they pay the practical cost of doing so; if they do not, they pay a moral cost, in the form of a reason for regret. While this double-bind does not always make requests wrongful, I show how it can be intrusive, manipulative, and coercive. Offers do not create the same double-bind, I argue, because offers do not as such create reasons.

I turn, finally, to a foundational question: How can the power to create a reason by requesting be explained *by its value*? It does not in general follow from the fact that it would be good for us to possess an ability that we do possess it: it would be good if we could fly like birds, but we can’t. Hence there is a worry that my account of requests, along with several prominent value-based explanations of requests, promises, and commands, rests on a fallacious form of wishful thinking. I argue that the power to create reasons by requesting can be justified by comparing its value *for each person*: a person possesses this power when, in general, it is better for *requesters* in these circumstances to *possess* the power than it would be for *requestees* to *not be subject* to the power. I generalize my account to show how commands and promises can be justified, and I argue my view does not rest on wishful thinking, but rather on a familiar and widely-accepted form of explanation in ethics, which justifies normative principles by considering what the world would be like for each of us, were those principles generally accepted.